

# Efficient RLVR (Data & Computation)

Hang Yang, Gio Song, Lisa Zhu

# Act Only When It Pays: Efficient Reinforcement Learning for LLM Reasoning via Selective Rollouts

By Hang Yang

# Background

## RL Powers LLM Reasoning

Reasoning models leverage **Chain-of-Thought** (CoT) for stronger reasoning (e.g., OpenAI o1, DeepSeek R1)

**Key driver:** Reinforcement learning (RL) enable iterative strategy refinement with PPO and GRPO

**Importantly,** at **rollout** stage, generating more prompts can **further enhance training**




## Rollout Scaling Benefits

higher-quality data

Stabilizes RL training

Improves model convergence

## The main Challenge - Computational Resources



# How to focus on sampling more valuable prompts?

## Existing Methods

Static Data Pruning




A new model to identify valuable data point  
**(No conclusive evidence of improving overall efficiency)**

Dynamic Sampling




The value of a data point varies across models  
and stages (**non-adaptive**)

oversampling and filter out uninformative data only  
after rollout (**additional rollout cost**)



Online data selection



Low computational overheads

From the observation and analysis of **GRPO** to propose a new algorithm **GRESO**

# Group Relative Policy Optimization (GRPO)

## Objective function

$$\mathcal{J}_{GRPO}(\theta) = \mathbb{E}[q \sim P(Q), \{o_i\}_{i=1}^G \sim \pi_{\theta_{old}}(O|q)]$$

$$\frac{1}{G} \sum_{i=1}^G \left( \min\left(\frac{\pi_\theta(o_i|q)}{\pi_{\theta_{old}}(o_i|q)}, 1-\epsilon, 1+\epsilon\right) A_i \right) - \beta \mathbb{D}_{KL}(\pi_\theta || \pi_{ref})$$

$$A_{i,t} = \frac{r_i - \text{mean}(\{R_i\}_{i=1}^G)}{\text{std}(\{R_i\}_{i=1}^G)}.$$

one prompt → a group of response corresponding with a group of rewards {r1, r2, r3,...,rG}  
Ai,t → advantage function to evaluate whether an example can provide learning signal

Prompts

Uninformative

Output1: 5

Output2: 5

...

OutputG: 5

informative

Output1: 1

Output2: 5

...

OutputG: 23

High Variance




Effective Prompts

# GRPO Observations

## Observation 1

Effective Prompts Ratio keeps **decreasing** as the training proceeds



Varying EPR hurt **training stability** and **final model performance**

Zero-Variance prompts

80%

Identifying **Priors** to Rollout




5 times Rollouts

Maintain batch size

## Observation 2

The information value of a prompt is continuous and predictable over time



$P(\text{Previous} \mid \text{Current})$ : 90% ~

$P(\text{Current} \mid \text{Previous})$  ~ 80%

in most cases it remains **consistent** (zero-variance stays zero-variance), but a small portion may **transition**.

Retain **Potentially Valuable** Prompts

# Algorithm GRPO with Efficient Selective Rollout (GRESO)

Identifying **Priorly** : formalize the problem of zero-variance prompt detection

$$T_i = (e_{i,1}, R_{i,1}), \dots, (e_{i,n}, R_{i,n}) \quad R_{i,1} = \{r_{i,1}^{(k)}\}_{k=1}^G$$

$e_{i,j}$  denotes the epoch number (example  $x_i$  and j-th sampling)

$R_{i,1}$  represents the set of response rewards


To **predict** whether  $x_i$  is a zero-variance prompt

**Probabilistic Pre-rollout Prompt Filtering :**

$$p_f(x_i) = 1 - p_e^{z_i},$$

$$z_i = \max \left\{ k \in [0, n] \left| \prod_{j=n-k+1}^n \mathbb{I}_{i,j} = 1 \right. \right\},$$

$$\mathbb{I}_{i,j} = \begin{cases} 1, & \text{if all rewards in } R_{i,j} \text{ are identical,} \\ 0, & \text{otherwise,} \end{cases}$$



[1, 1, 0, 1]



$Z_i = 2$

# Algorithm GRPO with Efficient Selective Rollout (GRESO)

## Probabilistic Pre-rollout Prompt Filtering :

$$p_f(x_i) = 1 - p_e^{z_i},$$

$p_e$  denotes base exploration probability ( $p_e \uparrow \quad p_f \downarrow$ )  
 $p_f$  denotes probability of Pre-rollout Prompt Filtering

```

1  $\mathcal{B} \leftarrow \emptyset; B_r \leftarrow B_r^{\text{default}}; n_{\text{easy}}, n_{\text{hard}}, n_{\text{total}} \leftarrow 0, 0, 0;$ 
2 /* Rollout Stage.
3 repeat
4    $\{x_i\}_{i=1}^{B_r} \leftarrow$  Sample prompts from  $\mathcal{D}$  and filter with Eq. 3 until batch size =  $B_r$ ;
5    $\{x_i, r_i\}_{i=1}^{B_r \times G} \leftarrow$  Rollout generation on  $\{x_i\}_{i=1}^{B_r}$ ;
6    $\{x_i, r_i\}_{i=1}^{B_r \times G} \leftarrow$  filter out zero-var prompt in  $\{x_i, r_i\}_{i=1}^{B_r \times G}$ ;
7    $n_{\text{easy}} \leftarrow n_{\text{easy}} +$  filtered easy zero-var prompt count;
8    $n_{\text{hard}} \leftarrow n_{\text{hard}} +$  filtered hard zero-var prompt count;
9    $n_{\text{total}} \leftarrow n_{\text{total}} + B_r$ ;
10   $\mathcal{B} \leftarrow \mathcal{B} \cup \{x_i, r_i\}_{i=1}^{B_r \times G}$ ;
11  /* Adaptive rollout batch size.
12   $B_r \leftarrow \min(B_r^{\text{default}}, \text{Adaptive rollout batch size calculated by Eq. 6})$ ;
13 until  $|\mathcal{B}| \geq B_t$ ;
14 /* Adjust Base Exploration Probability.
15 if  $n_{\text{easy}}/n_{\text{total}} \geq \alpha_{\text{easy}}$  then  $p_{\text{easy}} \leftarrow p_{\text{easy}} - \Delta p$  ;
16 else  $p_{\text{easy}} \leftarrow p_{\text{easy}} + \Delta p$  ;
17 if  $n_{\text{hard}}/n_{\text{total}} \geq \alpha_{\text{hard}}$  then  $p_{\text{hard}} \leftarrow p_{\text{hard}} - \Delta p$  ;
18 else  $p_{\text{hard}} \leftarrow p_{\text{hard}} + \Delta p$  ;
19 /* GRPO Training.
20  $\mathcal{B} \leftarrow$  select  $B_t$  examples from  $\mathcal{B}$ ;
21 Update actor model with GRPO on  $\mathcal{B}$ ;
```

Dynamically adjusting  $p_e$  and Batchsize



$n_{\text{filtered}}/n_{\text{total}} > \alpha \rightarrow p_e \downarrow$

More probability to filter zero-variance

$$B_r = \min \left( B_r^{\text{default}}, \beta \frac{B_\Delta}{(1 - \alpha)} \right)$$

Dynamical batchsize  $\rightarrow$  no extra waste

# Experiment

## End-to-end Efficiency Comparison





Dataset	Method	Math500	AIME24	AMC	Gaokao	Miner.	Olymp.	Avg.	# Rollout
<i>Qwen2.5-Math-1.5B</i>									
DM	DS	77.3	16.7	61.7	64.2	31.8	38.7	48.4	7.6M
	GRESO	76.6	15.0	61.4	66.2	33.3	38.5	48.5	<b>3.3M</b>
OR1	DS	77.1	16.7	50.3	65.5	30.9	39.7	46.7	3.8M
	GRESO	76.1	20.0	50.6	65.1	30.0	39.2	46.8	<b>1.6M</b>
<i>DeepSeek-R1-Distill-Qwen-1.5B</i>									
DM	DS	87.9	36.7	71.7	78.7	35.3	55.9	61.0	2.4M
	GRESO	87.7	36.7	71.1	78.4	33.9	55.1	60.5	<b>1.6M</b>
OR1	DS	84.8	25.0	68.4	74.0	34.1	54.2	56.7	2.4M
	GRESO	85.9	26.7	66.9	75.2	33.6	55.5	57.3	<b>1.5M</b>
<i>Qwen2.5-Math-7B</i>									
DM	DS	82.9	34.2	79.2	71.7	35.4	43.6	57.8	13.1M
	GRESO	82.2	32.5	80.7	70.2	35.4	44.1	57.5	<b>6.3M</b>
OR1	DS	82.9	34.2	63.1	67.3	34.9	46.3	54.8	11.4M
	GRESO	82.3	35.0	64.5	66.8	36.5	45.7	55.1	<b>3.4M</b>

Method	Training	Other	Rollout	Total
<i>Qwen2.5-Math-1.5B</i>				
DS	8.1	3.6	41.0 (1.0×)	52.6 (1.0×)
GRESO	8.9	3.9	<b>25.2 (1.6×)</b>	<b>37.9 (1.4×)</b>
<i>DeepSeek-R1-Distill-Qwen-1.5B</i>				
DS	6.1	3.3	92.4 (1.0×)	101.9 (1.0×)
GRESO	6.8	4.0	<b>62.0 (1.5×)</b>	<b>72.7 (1.4×)</b>
<i>Qwen2.5-Math-7B</i>				
DS	16.1	6.1	155.9 (1.0×)	178.0 (1.0×)
GRESO	16.6	6.3	<b>65.5 (2.4×)</b>	<b>88.3 (2.0×)</b>

No performance drop with  
up to **3.35×** fewer rollouts and up to **2.4x** wall-clock time speed-up

# Experiment

## Analysis and Ablation Study






**DS:** Filters zero-variance prompts after rollout, but effective ratio drops and costs rise

**GRESO:** Skips zero-variance prompts before rollout, keeping >70% effective ratio and lower cost

# Experiment

## Dynamics of self-adjustable base exploration probabilities.



GRESO adaptively adjusts exploration probabilities without manual tuning  
As the model improves, Pe increases to explore harder examples

# Conclusion

## Key Contribution

GRESO : Act only when it pays, a novel algorithm to optimize rollout selection

3.35x fewer  
rollouts

2.4x rollout  
Speed up

2.0x overall  
training  
Speed up

## Future Prospects

Extending selective rollouts to broader domains and more sophisticated data selection

# Beyond the 80/20 Rule: High-Entropy Minority Tokens Drive Effective Reinforcement Learning for LLM Reasoning

By Gio Song

# Background

## Why Token-Level Analysis in RLVR Matters

- Reinforcement Learning for Verifiable Reasoning (RLVR) has become the **standard alignment method** for LLMs. But it shows only *moderate* gains
- Most prior work focuses on:
  - Algorithmic innovation (e.g., DAPO)
  - Task adaptation beyond math (e.g., Absolute Zero)
  - Empirical tricks (e.g., One-shot training)
- ! Missing: analysis of **how specific tokens** contribute to performance

# Why This Paper?

## So What Are We Missing in RLVR?

- Prior work treats all tokens equally during training
- But not all tokens are equally important in reasoning!
- Question: Can we identify and optimize the *right tokens*?

### Quote for emphasis:

- “High-entropy tokens may decide reasoning *paths*, not just language *forms*.”
- Studying tokens, in fact, means studying the conditional probability distribution of the next token output by an LLM.


# Key insights

Token Type	Entropy	Role in Output
Low-entropy	Very stable	Fills in predictable structure (e.g., math formulas, code)
High-entropy	Uncertain	Drives reasoning direction; controls "forks" in logic

Example:

In decimal,  $1+1=2$ . But how does that translate to base 2? Well, in binary [...]

- Blue tokens = low-entropy; ● red tokens = high-entropy (forking tokens)



# Further Discoveries

- **Slightly increasing entropy of high-entropy tokens** improves performance
- RLVR primarily **adjusts the entropy of high-entropy tokens**, while low-entropy tokens remain largely unchanged


# Main Experiment & Ablation Experiment

Based on earlier findings, the authors hypothesize that:

- **Optimizing the conditional distributions of low-entropy tokens is unnecessary.**
- Instead, **only high-entropy tokens** ( $\approx 20\%$  of all tokens) need targeted gradient updates to replicate most of the RL benefits.

The authors also **tune the proportion** of tokens to treat as “high-entropy” and find:

- **20% is optimal** for balancing performance and gradient efficiency.



# Preliminaries

## 1. Token Entropy

Token entropy is based on the conditional probability distribution **over the vocabulary** at each step, not the specific token identity.

$$H_t = - \sum_{j=1}^V p_{t,j} \log p_{t,j}, \quad \text{where } p_t = \text{Softmax} \left( \frac{z_t}{T} \right)$$

## 2. DAPO – Dynamic sAmpling Policy Optimization

- DAPO selects **partially correct** prompts for training.
- Encourages learning from **useful but imperfect** trajectories.
- Advantage estimation ensures training focuses on **relatively better samples**.

$$\mathcal{J}_{\text{DAPO}}(\theta) = \mathbb{E} \left[ \frac{1}{\sum_{i=1}^G |o^i|} \sum_{i=1}^G \sum_{t=1}^{|o^i|} \min \left( r_t^i(\theta) \hat{A}_t^i, \text{clip}(r_t^i(\theta), 1 - \epsilon_{\text{low}}, 1 + \epsilon_{\text{high}}) \hat{A}_t^i \right) \right]$$

# Pre--Experiment

### 3.1 Token Entropy in Chain-of-Thought (CoT)

- **Goal:** Analyze entropy distributions in CoT outputs

## Key Analysis:

- **Token Entropy Distribution:**


- Only **20% of tokens** have entropy > **0.672**
  - Most tokens are low-entropy — structural
  - High-entropy tokens are rare, but impactful

- Word Cloud Visualization


## Conclusion:

High-entropy tokens play a **decisive role in branching logic**


**They are termed “forking tokens”**



(a) Distribution of token entropy



(b) Frequent tokens with the highest average entropy



(c) Frequent tokens with the lowest average entropy

# Entropy Intervention Experiment




Figure 3: Average scores of AIME 2024 and AIME 2025. Red curve varying  $T_{high}$  with  $T_{low} = 1$ . Blue curve varying  $T_{low}$  with  $T_{high} = 1$ .

## Method:

- Define threshold:  $H_{threshold}=0.672$
- Use adaptive temperature scaling:

$$T'_t = \begin{cases} T_{high} & \text{if } H_t > H_{threshold} \\ T_{low} & \text{otherwise} \end{cases}$$

- Test two conditions:
  - Fix  $T_{low}=1$ , vary  $T_{high}$  (**Red Curve**)
  - Fix  $T_{high}=1$ , vary  $T_{low}$  (**Blue Curve**)

## Insight:

Selectively **increasing entropy at forking tokens** improves Reasoning

This **mirrors the effect of RL training**, where entropy change is concentrated at decision-critical points

# Pre--Experiment

## 3.2: RLVR Retains and Strengthens Entropy Patterns of Base Models

### 1) RLVR Retains Entropy Structure of the Base Model

Compare the **top 20% high-entropy tokens** between:

- **Base model**
- **Intermediate RLVR models**
- **Final RLVR model**



86% of high-entropy tokens remain consistent

Table 1: The progression of the overlap ratio in the positions of the top 20% high-entropy tokens, comparing the base model (i.e., step 0) with the model after RLVR training (i.e., step 1360).

Compared w/	Step 0	Step 16	Step 112	Step 160	Step 480	Step 800	Step 864	Step 840	Step 1280	Step 1360
Base Model	100%	98.92%	98.70%	93.04%	93.02%	93.03%	87.45%	87.22%	87.09%	86.67%
RLVR Model	86.67%	86.71%	86.83%	90.64%	90.65%	90.64%	96.61%	97.07%	97.34%	100%

# Pre--Experiment

## 3.2: RLVR Retains and Strengthens Entropy Patterns of Base Models

### 2) RLVR Selective Entropy Adjustment:

- Tokens grouped by **5% entropy percentile intervals** (from low to high)
- Compute **average entropy change** after RLVR for each group




Figure 4: Average entropy change after RLVR within each 5% entropy percentile range of the base model.  $x\%$  percentile means that  $x\%$  of the tokens in the dataset have entropy values less than or equal to this value. It is worth noting that the Y-axis is presented on a *log scale*. Tokens with higher initial entropy tend to experience greater entropy increases after RLVR.

RLVR keeps the original token distribution structure intact but **selectively increases entropy for a small set** (top 20%) of tokens This sets the foundation for training **only high-entropy tokens** in later sections.

# Main--Experiment

Adapted DAPO objective for only **high-entropy tokens**:


$$\mathcal{J}_{\text{HighEnt}}^B(\theta) = \mathbb{E} \left[ \dots \mathbb{I}(H_t^i \geq \tau_p^B) \min(r_t^i(\theta) \hat{A}_t^i, \text{clip}(\cdot)) \right]$$

- Only tokens with entropy  $\geq$  top-p threshold are used
- This means **RL updates only the most informative tokens**

Table 2: Comparison between *vanilla DAPO* using *all tokens* and *DAPO* using *only the top 20% high-entropy tokens* (i.e. *forking tokens*) in policy gradient loss, evaluated on the *Qwen3-32B*, *Qwen3-14B* and *Qwen3-8B* base models. "Acc@16" and "Len@16" denotes the average accuracy and response length over 16 evaluations per benchmark, respectively.

Benchmark	DAPO w/ All Tokens		DAPO w/ Forking Tokens		Improvement	
	Acc@16	Len@16	Acc@16	Len@16	Acc@16	Len@16
RLVR from the Qwen3-32B Base Model						
AIME'24	55.83	9644.15	<b>63.54</b>	12197.54	+ <b>7.71</b>	+2553.39
AIME'25	45.63	9037.48	<b>56.67</b>	11842.25	+ <b>11.04</b>	+2804.77
AMC'23	91.88	5285.03	<b>94.22</b>	5896.47	+ <b>2.34</b>	+611.44
MATH500	94.36	2853.51	<b>94.88</b>	3366.01	+ <b>0.52</b>	+512.5
Minerva	45.70	2675.28	<b>45.82</b>	2759.88	+ <b>0.12</b>	+84.6
Olympiad	66.16	5597.37	<b>69.02</b>	7300.01	+ <b>2.86</b>	+1702.64
<b>Average</b>	<b>66.59</b>	5848.80	<b>70.69</b>	7227.03	+ <b>4.10</b>	+1378.22
RLVR from the Qwen3-14B Base Model						
AIME'24	45.21	7945.15	<b>50.42</b>	11814.36	+ <b>5.21</b>	+3869.21
AIME'25	38.13	7056.98	<b>42.92</b>	12060.48	+ <b>4.79</b>	+5003.5
AMC'23	89.53	4509.37	<b>91.56</b>	7095.13	+ <b>2.03</b>	+2585.76
MATH500	92.23	2348.22	<b>93.59</b>	3970.10	+ <b>1.37</b>	+1621.88
Minerva	42.16	2011.16	<b>43.20</b>	2959.32	+ <b>1.03</b>	+948.16
Olympiad	61.14	4642.07	<b>64.62</b>	7871.25	+ <b>3.48</b>	+3229.18
<b>Average</b>	<b>61.40</b>	4752.16	<b>64.39</b>	7628.44	+ <b>2.99</b>	+2876.28
RLVR from the Qwen3-8B Base Model						
AIME'24	33.33	6884.89	<b>34.58</b>	9494.29	+ <b>1.25</b>	+2609.40
AIME'25	25.42	5915.91	<b>26.25</b>	8120.20	+ <b>0.83</b>	+2204.29
AMC'23	<b>77.81</b>	3967.91	77.19	5450.62	- <b>0.625</b>	+1482.71
MATH500	89.24	2059.00	<b>89.70</b>	2672.91	+ <b>0.46</b>	+613.91
Minerva	39.77	1450.68	<b>40.26</b>	2068.41	+ <b>0.48</b>	+617.73
Olympiad	56.67	3853.55	<b>57.43</b>	5241.54	+ <b>0.76</b>	+1387.99
<b>Average</b>	<b>53.71</b>	4021.99	<b>54.23</b>	5508.00	+ <b>0.53</b>	+1486.01

Reinforcement learning performance boost is largely driven by forking tokens



# Further--Experiment

1. Varying  $\rho$  (proportion of high-entropy tokens)

2. Model Size Impact







Figure 8: Scaling trend of DAPO using only forking tokens (i.e., top 20% of high-entropy tokens) in policy gradient loss. These results suggest that concentrating exclusively on forking tokens in the policy gradient loss may yield greater benefits in larger reasoning models.

**Smaller subset of tokens (high entropy) can drive stronger performance, reducing cost while increasing quality.**


--foundational claim of the article




(a) AIME'24 scores trained from Qwen3-8B base.




(b) Overall entropy trained from Qwen3-8B base.




(c) AIME'24 scores trained from Qwen3-14B base.



(d) Overall entropy trained from Qwen3-14B base.








(e) AIME'24 scores trained from Qwen3-32B base.



(f) Overall entropy trained from Qwen3-32B base.

# Analysis

Aspect	Finding
Cross-task generalization	High-entropy token updates <b>improve transfer</b> (math → code)
Long-context reasoning	Training with forking tokens supports longer outputs and deeper logic
Portability to smaller models	Works well even under <b>low-compute, small-model cold-start</b> scenarios. <b>model-agnostic</b>



(a) AIME'24 scores trained from Llama-3.1-8B after cold-start.

(b) Response length trained from Llama-3.1-8B after cold-start.

# Discussion, Conclusion & Limitations

## Discussion & Conclusions

- Why High-Entropy Tokens Matter in RL
- LLM CoT and Token Entropy
- Why RLVR Works



## Develop better RLVR algorithms

- Supervised fine-tuning (SFT)
- Distillation
- Inference pipelines
- Multi-modal training

## Limitations & Further Improvement

- Mainly on **Qwen models**.
- Dataset limited to mathematical reasoning.
- Results are experiment-specific.



# Spurious Rewards: Rethinking Training Signals in RLVR

Lisa Zhu, Hang Yang, Gio Song

# Core Idea & Findings


- Reinforcement Learning with Verifiable Rewards (RLVR) improves reasoning in LLMs
- Surprisingly, it works even with spurious rewards
  - Random, wrong, or irrelevant
- Qwen2.5-Math-7B
  - Random rewards: 21.4%
  - Wrong label: +24.1%
- Performance gains nearly match ground truth training

# Additional Insights

- Model differences
  - Strong gains for Qwen2.5-Math
  - Little or negative effect on Llama3 & OLMo2
- Code reasoning (thinking in code without actual code execution):
  - Distinctive behavior for Qwen2.5-Math
  - Becomes more frequent after RLVR
  - From 65% → 90%
- Implication
  - RLVR surfacing latent abilities from pretraining
    - Not reward signal itself

# Experiment & Results I

- Goal: Test if RLVR still improves reasoning with weaker or spurious rewards instead of ground truth
- Method:
  - Base model: Qwen2.5-Math
  - Training: GRPO algorithm, DeepScaleR dataset
    - GRPO finetune base model
    - DeepScaleR trained with spurious binary (0-1) reward functions
- Investigate the limits of how little supervision is needed for RLVR training



# Experiment & Results II

- Types of rewards tested
  - Standard to Weak to Spurious
  - Ground Truth → Majority Vote  
→ Format → Random  
→ Incorrect
    - Ground Truth: Correct answers only
    - Incorrect: Deliberately reward wrong answers from pseudo-labeling
- Results
  - All reward types have significant math gains within 1st 50 steps
  - Smaller model also improves, but more slowly
  - Takeaway: RLVR boosts performance even with spurious signals
    - Elicit **latent abilities** from pretraining

# CROSS- Model Analysis

- Goal: Test if spurious-reward gains generalize across models
- Models: Qwen2.5-Math vs. OLMo2 vs. Llama3
- Findings:
  - Qwen2.5-Math: large gains even with spurious rewards
  - OLMo2 & Llama3: minimal or negative gains
- Why Qwen
  - Exhibits strong code reasoning (i.e. writes math steps in Python)
  - Accuracy: 61% with code vs. 28% without
- Takeaway: Spurious rewards amplify latent code-reasoning abilities in Qwen, not transferable to other model

# What makes Qwen so SPECIAL?--Code Reasoning

- **Observation:** The Qwen2.5-Math model frequently generates Python code as a method of reasoning
- **More structured and accurate solutions.**
- **Evidence: 65%+ code reasoning frequency.** After RLVR (even with random or incorrect rewards), this frequency rises above **90%**.
- **Why this matters:**  
Code reasoning leads to significantly **higher accuracy (60.9%)**. Only **35.0%** on responses with natural language reasoning.
- **Other models (e.g., OLMo, LLaMA)** either do not use code (No-Code) or use it ineffectively (Bad-Code), and hence don't benefit similarly from RLVR.

## MATH Question:

What is the distance, in units, between the points  $(2, -6)$  and  $(-4, 3)$ ? Express your answer in simplest radical form.

## Qwen2.5-Math-7B Solution (correct):

To find the distance between two points  $(x_1, y_1)$  and  $(x_2, y_2)$  in a Cartesian plane...  
Let's break this down step-by-step and compute the result using Python.

```
1 import math
2 ...
3 # Calculate the distance using the distance formula
4 distance = math.sqrt(dx**2 + dy**2)
5 print(distance)
```

output: 10.816653826391969


...

Thus, the final answer is:  $3\sqrt{13}$

Model	Qwen2.5-Math-7B	Qwen2.5-Math-1.5B	Qwen2.5-7B	OLMo2-7B-SFT
Code Frequency	65.0	53.6	92.2	98.0
Acc. w/ Code	60.9	52.6	39.9	21.0
Acc. w/ Lang	35.0	17.2	61.5	40.0

# RLVR with Spurious Rewards Amplifies Pretrained Reasoning Strategies

- Why do spurious rewards work?
- **Evidence:** Code Reasoning Frequency Strongly Correlates with Accuracy
- **Before RLVR:** Qwen2.5-Math-7B uses code reasoning in 65% of outputs.
- **After RLVR:** rises to 90–95%, and accuracy increases alongside.
- **Random reward** leads to slower increase but eventually hits 95.6% code reasoning rate.
- **True label reward** causes an initial spike in code usage, but this later **declines** as the model learns to solve more via natural language.



# RLVR with Spurious Rewards Amplifies Pretrained Reasoning Strategies



The authors examine performance shifts across 4 reasoning transition patterns:

<b>Code→Code</b>	Code reasoning before and after training
<b>Code→Lang</b>	Switch from code to language reasoning
<b>Lang→Code</b>	Switch from language to code reasoning
<b>Lang→Lang</b>	Natural language reasoning both before and after

Two main metrics were tracked:

- Subset **frequency** (how often that strategy occurred)
- Subset **accuracy** (how correct it was)

# RLVR with Spurious Rewards Amplifies Pretrained Reasoning Strategies



## Findings from Strategy Shift Analysis:

- Under **spurious and weak rewards**, Qwen2.5-Math-7B tends to:
  - Maintain code reasoning if it already used it. (Code→Lang)
  - Switch from language to code reasoning** (Lang→Code) in most other cases.
- True reward** does not cause the same shift

Other models behave differently:

- Qwen2.5-7B** sees a **decline in code reasoning** under correct/majority/incorrect rewards
- OLMo2-7B-SFT** also shows **decreased code use** under valid reward signals.
- LLaMA and other No-Code models** show no meaningful change in strategy.

# Analysis

Table 2: Partial contribution to the overall performance gain averaged over rewards that successfully steered the model’s reasoning strategy (Figure 6).

Model	Qwen2.5-Math-7B	Qwen2.5-Math-1.5B	Qwen2.5-7B
Avg. Total Gain	↑ 23.5%	↑ 28.5%	↑ 30.6%
$C_{\text{Code} \rightarrow \text{Code}}$	11.6%	2.8%	0.2%
$C_{\text{Code} \rightarrow \text{Lang}}$	8.6%	2.0%	<b>93.9%</b>
$C_{\text{Lang} \rightarrow \text{Code}}$	<b>58.3%</b>	<b>78.7%</b>	0.0%
$C_{\text{Lang} \rightarrow \text{Lang}}$	21.4%	16.5%	5.9%

- **Qwen-Math models improve by switching into their strength (code reasoning).**
- **Other models improve by abandoning inefficient strategies, like code reasoning, in favor of simpler text reasoning.**
- For Qwen2.5-Math, the performance gains from spurious reward **do not reflect new skill acquisition**, but rather **the amplification of a previously learned, effective strategy** (code reasoning).
- **RLVR, particularly with non-informative or even misleading reward signals, can still work extremely well — if and only if the underlying model has already internalized useful reasoning strategies during pretraining.**

# Interventions on code reasoning

## Impact of Increased Code Reasoning on Performance

(1) Prompting (Answer begin with “let’s solve this using python”)

Model	Original	Prompting	Abs. Diff.
Qwen2.5-Math-1.5B	36.2%	60.4%	+24.2%
Qwen2.5-Math-7B	49.4%	64.4%	+15.0%
Qwen2.5-1.5B	3.0%	13.0%	+10.0%
Qwen2.5-7B	41.6%	22.2%	-19.4%
Llama3.2-3B-Instruct	36.8%	8.2%	-28.6%
Llama3.1-8B-Instruct	36.8%	15.2%	-21.6%
OLMo2-7B	9.0%	7.8%	-1.2%
OLMo2-7B-SFT	21.4%	18.6%	-2.8%


Legend:

- Qwen-Math-7B
- Qwen-Math-1.5B
- Qwen-7B
- Qwen-1.5B
- Olmo2-7B-SFT
- Olmo2-7B
- Llama3.1-8B
- Llama3.2-3B
- Llama3.1-8B-Instruct
- Llama3.2-3B-Instruct

Qwen model : 

Llama, OLMo : 

(2) RLVR(Assign a positive rewards only answer contain “python”)





Qwen2.5-Math-7B model generated code reasoning in its' answer >99% just 20 training steps

# Inhibiting code reasoning during RLVR with spurious rewards

Reward a response if and only if:

(1) spurious reward condition (**original**) (2) no string "python" (**compound**)



(a) Qwen2.5-Math-7B

Qwen math model : (1) format reward ↓ (2) Incorrect reward (AMC ↓)

(3) Ground truth Performance improvement ≠ sole code reasoning frequency

Bad code model : Compoud rewards > Original (downweight suboptimal model behavior)

# Curious cases : Training Signals from Incorrect Rewards and Random Rewards

## Hypothesis: Incorrect Rewards → Reasoning

- (1) many incorrect labels remain close to ground truth values ( positive reinforcement)
- (2) incorrect labels may function like format reward (some degree of correct )

## Random Rewards → Reasoning

Hypothesis from someone : most rewarded answers are correct (X)


Rewarded response : correct > incorrect  Penalized response : correct > incorrect

Normalization of reward in GRPO  Random rewards ≠ bias toward correct answers

Why random rewards worked ?

# Why random rewards worked?

## Experiment 1 : Random rewards with varying probabilities




Except for  $\gamma = 0$ ,  
 $\gamma$  do not affect the final performance

## Experiment 2: Clipping function enabled Vs disabled

$$\mathcal{J}_{GRPO}(\theta) = \mathbb{E}[q \sim P(Q), \{o_i\}_{i=1}^G \sim \pi_{\theta_{old}}(O|q)]$$

$$\frac{1}{G} \sum_{i=1}^G \left( \min\left(\frac{\pi_\theta(o_i|q)}{\pi_{\theta_{old}}(o_i|q)}, 1-\epsilon, 1+\epsilon\right) A_i, \text{clip}\left(\frac{\pi_\theta(o_i|q)}{\pi_{\theta_{old}}(o_i|q)}, 1-\epsilon, 1+\epsilon\right) A_i \right) - \beta \mathbb{D}_{KL}(\pi_\theta || \pi_{ref})$$



- (a) Performance on MATH-500
  - (b) Frequency of Code Reasoning
  - (1) directly turning off the clipping term
  - (2) adjusting training and rollout batch sizes ( $\pi_\theta = \pi_{old}$ )
- Clipping: ~21% performance gain

Optimizing algorithm's bias toward exploiting priors learned  
during pretraining (Amplify penalties, Regulate rewards)

# Conclusion

## Summary

- (1) RLVR with spurious rewards (random, incorrect, format-only) improves Qwen2.5-Math by amplifying pre-existing code reasoning patterns rather than teaching new skills.
- (2) Code reasoning frequency increases from 65% to 90%+ during training, directly correlating with performance gains across all reward types.
- (3) Model-dependent effects — spurious rewards work for Qwen families but consistently fail for Llama and OLMo models

## Key Implications

- (1) Pretraining determines outcomes — RLVR effectiveness depends on what reasoning patterns already exist in the base model.
- (2) Spurious signals can work — when they trigger beneficial pre-trained behaviors like code reasoning capabilities.

# R-Zero: Self-Evolving LLM from Zero Data

By Lisa Zhu

# Motivation

- LLMs need huge amounts of **human-curated data and labels** for fine-tuning
- **Costly, slow, and limits scalability** toward true self-evolving AI
- Existing “label-free” methods still rely on **pre-existing tasks or external verification**
- **R-Zero:** Fully **autonomous** framework
  - LLMs generates **its own training data** from scratch

# Preliminaries


## Group Relative Policy Optimization(GRPO)

- **Reinforcement Learning algorithm** for fine-tuning LLMs
- ~~Separate value function~~ Compares responses within the same group
- Uses **z-score normalization** of rewards: each answer is judged relative to others
- Encourages better responses while preventing large policy drift

## Reinforcement Learning with Verifiable Rewards (RLVR)

- Paradigm for fine-tuning models
- Applies when response quality can be objectively checked
- Uses rule-based verifier
  - Reward = 1 if correct, 0 if wrong
- Foundation for training the Solver in R-Zero

# Methodology Overview



- R-Zero = Challenger + Solver, initialized from the same LLM.
- Works in an iterative loop:
  - Challenger generates synthetic questions via GRPO.
  - Solver trains on these questions with pseudo-labels.
- Self-supervised: no human labels required.
- Goal: Challenger and Solver **co-evolve**, making Solver increasingly stronger

# Challenger & Solver Training

## Challenger ( $Q\theta$ )

- Generates **challenging questions** via GRPO.
- Guided by reward signals (uncertainty, penalties).
- Goal: push Solver to face progressively harder tasks

## Solver ( $S\phi$ )

- Fine-tuned on Challenger's filtered question set.
- Uses GRPO with a **verifiable reward**:  
$$r_j = \begin{cases} 1, & \text{if } x_j \text{ is identical to the pseudo-label } \tilde{y}_i, \\ 0, & \text{otherwise.} \end{cases}$$
- Learns to correctly answer increasingly difficult questions

# Reward Function – Uncertainty Reward

$$r_{\text{uncertainty}}(x; \phi) = 1 - 2 \left| \hat{p}(x; S_\phi) - \frac{1}{2} \right|$$

- Encourages **questions with mid-level difficulty.**
- Solver's accuracy on question  $x$ :  
$$\hat{p}(x; S_\phi) = \frac{1}{m} \sum_{j=1}^m \mathbb{1}\{y_j = \tilde{y}(x)\}$$
- Maximized when Solver accuracy  $\approx 50\%$ , forcing learning on “frontier” problems

# Repetition & Format Penalties

- **Repetition Penalty**

- Prevents generating near-duplicate questions.
- Uses **BLEU score similarity**; larger clusters → larger penalty.
- Formula:

$$r_{\text{rep}}(x_i) = \lambda \frac{|C_k|}{B}$$

- **Format Check Penalty**

- Structural rule: question must be enclosed in <question> & </question>
- If not, reward = 0 and question is discarded

# Reward Function – Composite Reward

- Purpose: Combine signals from uncertainty and repetition to train Challenger effectively.

- Formula:

$$r_i = \max(0, r_{\text{uncertainty}}(x_i; \phi) - r_{\text{rep}}(x_i))$$

- Interpretation:

- Starts from **uncertainty reward** (challenging but solvable questions).
- Subtracts penalty if question is too similar to others.
- Ensures reward  $\geq 0$ , preventing negative reinforcement.

- **Takeaway:** Final reward signal balances *difficulty* with *diversity*



## Experiments Setup – Models & Training

- Models Tested
  - **Qwen3-4B / 8B** → scale within same family
  - **OctoThinker-3B / 8B** → different lineage (Llama-based)
  - Ensures evaluation across **two distinct architectures**
- Training Details
  - Candidate pool: **8,000 questions** per iteration
  - Solver samples 10 answers per question
  - Keep only mid-consistency tasks (**3–7 matched answers**)
  - **Rewards:** uncertainty (Solver confusion)

# Experiments Setup – Benchmarks

- Mathematical Reasoning
  - 7 Benchmarks: AMC, Minerva, MATH-500, GSM8K, OlympiadBench, AIME-2024, AIME-2025
  - Test correctness, complexity, and comprehensiveness
  - Metrics reported:
    - AMC & AIME: mean@123
    - Others: accuracy (greedy decoding)
- General Domain Reasoning
  - **MMLU-Pro**: Harder multi-task questions (language model capabilities)
  - **SuperGPQA**: Graduate-level reasoning across 285 disciplines
  - **BBEH**: More difficult BIG-Bench tasks for complex reasoning

# Math Reasoning Results

- Findings
  - **Consistent gains across all models** (Qwen3 & OctoThinker families)
  - **Qwen3-8B:** +5.51 points (49.18 → 54.69 after 3 iterations)
  - **OctoThinker-3B:** +2.68 points (26.64 → 29.32)
  - Larger models improve more, but smaller ones still benefit
  - Takeaway: R-Zero is **effective & model-agnostic**, boosting performance across scales and architectures

Scores improve with each iteration; first iteration already gives a strong boost, showing RL-trained Challenger is critical

Model Name	AVG	AMC	Minerva	MATH	GSM8K	Olympiad	AIME25	AIME24
<i>Qwen3-4B-Base</i>								
Base Model	42.58	45.70	38.24	68.20	87.79	41.04	6.15	10.94
Base Challenger	44.36	45.00	45.22	72.80	87.87	41.19	7.29	11.15
R-Zero (Iter 1)	48.06	51.56	51.47	78.60	91.28	43.85	9.17	10.52
R-Zero (Iter 2)	48.44	52.50	51.47	79.80	91.66	44.30	4.27	15.10
R-Zero (Iter 3)	49.07	57.27	52.94	79.60	92.12	44.59	4.27	12.71
<i>Qwen3-8B-Base</i>								
Base Model	49.18	51.95	50.00	78.00	89.08	44.74	16.67	13.85
Base Challenger	51.87	60.70	57.72	81.60	92.56	46.44	13.44	10.62
R-Zero (Iter 1)	53.39	61.56	59.93	82.00	93.71	48.00	14.17	14.37
R-Zero (Iter 2)	53.84	61.56	59.93	82.00	93.93	48.30	17.60	13.54
R-Zero (Iter 3)	54.69	61.67	60.66	82.00	94.09	48.89	19.17	16.35
<i>OctoThinker-3B</i>								
Base Model	26.64	17.19	24.26	55.00	73.69	16.15	0.21	0.00
Base Challenger	27.51	20.19	24.63	54.60	74.98	15.70	0.10	2.40
R-Zero (Iter 1)	27.76	20.39	25.74	54.60	75.51	16.30	0.10	1.67
R-Zero (Iter 2)	28.20	24.06	25.37	54.80	74.45	17.48	0.00	1.25
R-Zero (Iter 3)	29.32	27.03	27.57	54.20	74.98	18.22	3.23	0.00
<i>OctoThinker-8B</i>								
Base Model	36.41	32.11	41.91	65.20	86.96	26.52	1.56	0.62
Base Challenger	36.98	29.30	42.28	66.20	88.10	27.56	1.04	4.38
R-Zero (Iter 1)	37.80	32.97	45.22	65.60	86.96	28.44	1.98	3.44
R-Zero (Iter 2)	38.23	32.58	48.53	67.20	87.11	27.26	0.00	4.90
R-Zero (Iter 3)	38.52	34.03	48.22	68.80	87.19	27.56	0.42	3.44

# General Results Reasoning

Model Name	Overall AVG	MATH AVG	SuperGPQA	MMLU-Pro	BBEH
<i>Qwen3-4B-Base</i>					
Base Model	27.10	42.58	20.88	37.38	7.57
Base Challenger	30.83	44.36	24.77	47.59	6.59
R-Zero (Iter 1)	34.27	48.06	<b>27.92</b>	51.69	9.42
R-Zero (Iter 2)	<b>34.92</b>	48.44	27.72	<b>53.75</b>	9.76
R-Zero (Iter 3)	34.64	<b>49.07</b>	27.55	51.53	<b>10.42</b>
<i>Qwen3-8B-Base</i>					
Base Model	34.49	49.18	28.33	51.80	8.63
Base Challenger	36.43	51.87	30.12	54.14	9.60
R-Zero (Iter 1)	37.93	53.39	31.26	57.17	9.91
R-Zero (Iter 2)	38.45	53.84	<b>31.58</b>	58.20	10.20
R-Zero (Iter 3)	<b>38.73</b>	<b>54.69</b>	31.38	<b>58.23</b>	<b>10.60</b>
<i>OctoThinker-3B</i>					
Base Model	12.27	26.64	10.09	10.87	1.46
Base Challenger	14.41	27.51	11.19	14.53	<b>4.40</b>
R-Zero (Iter 1)	14.93	27.76	12.21	15.72	4.05
R-Zero (Iter 2)	15.11	28.20	12.43	16.08	3.74
R-Zero (Iter 3)	<b>15.67</b>	<b>29.32</b>	<b>12.44</b>	<b>16.71</b>	4.20
<i>OctoThinker-8B</i>					
Base Model	16.81	32.11	13.26	20.21	1.64
Base Challenger	25.08	36.41	16.99	41.46	5.46
R-Zero (Iter 1)	26.44	37.80	19.15	<b>42.05</b>	6.77
R-Zero (Iter 2)	26.77	38.23	19.27	41.34	8.25
R-Zero (Iter 3)	<b>26.88</b>	<b>38.52</b>	<b>19.82</b>	40.92	8.25

These gains are not domain-specific — they generalize beyond math and enhance core reasoning ability


- **Findings:**
  - R-Zero improves **all tested models** in general reasoning
  - **Qwen3-8B:** +3.81 points (34.49 → 38.73)
  - **OctoThinker-3B:** +3.65 points (12.27 → 15.67)
  - Iterative gains across 3 rounds, similar to math results
- **Takeaway:** R-Zero's math-based training transfers to general reasoning skills

# Analysis – Ablation Study

- Removing **RL-Challenger, Filtering, or Repetition Penalty** → sharp performance drop.
- Biggest loss: without RL-Challenger ( $-3.7$  math,  $-4.1$  general).
- **Takeaway:** Each module is essential; Challenger RL drives curriculum quality

Method	Math AVG	General AVG
R-Zero (full)	48.06	30.41
<i>Ablations</i>		
└ w/o RL-Challenger	44.36	26.32
└ w/o Rep. Penalty	45.76	27.56
└ w/o Filtering	47.35	24.26

# Analysis – Difficulty & Synergy



Performance of Evaluated Model (vs. Ground Truth)					
	Base Model	Solver (Iter 1)	Solver (Iter 2)	Solver (Iter 3)	Pseudo-Label Acc.
$D_{\text{Iter } 1}$	48.0	59.0	57.0	61.0	79.0%
$D_{\text{Iter } 2}$	52.5	53.0	51.5	53.5	69.0%
$D_{\text{Iter } 3}$	44.0	47.0	45.0	50.5	63.0%


- **Difficulty Evolution:** Challenger makes tasks harder each round, but pseudo-label accuracy falls (79% → 63%)
- **Synergy with Human Labels:** Adding labeled data after R-Zero training yields **+2.35 points** over supervised baseline
- **Takeaway:** R-Zero improves difficulty handling, and works even better when combined with human labels

# Analysis – Scaling & Design

- **Iteration Scaling:** Larger models delay collapse; small models degrade earlier.
- **Label Noise:** Collapse linked to declining pseudo-label accuracy (but not the sole factor).
- **Two-Model Design:** Separate Challenger & Solver sustains higher performance (49.07 vs 45.57 for Single-R-Zero).
- **Takeaway:** Bigger models and two-model design stabilize training, but collapse risk remains.

Iteration	R-Zero (ours)		Single-R-Zero	
	Performance	Pseudo-label Acc (%)	Performance	Pseudo-label Acc (%)
Iter 1	48.06	71.0	47.31	63.4
Iter 2	48.44	56.2	46.95	46.6
Iter 3	49.07	48.8	45.57	32.6
Iter 4	46.52	42.2	43.89	33.8

Iteration	Model Size		
	0.6B	1.7B	4B
Iter 1	70.6	69.4	71.0
Iter 2	53.4	55.2	56.2
Iter 3	50.8	52.2	48.8
Iter 4	44.0	45.2	42.2



# Conclusion

- Contribution: R-Zero is the first framework to evolve reasoning LLMs with no external data
- Impact: Moves toward more autonomous & scalable AI training
- Limitations
  - Works best in domains with objectively verifiable answers (math)
  - Remains challenge in open-ended domains
- Future Directions
  - Improve label quality
  - Extend to broader reasoning
  - Prevent long-term collapse

Thank you!